

# Classifying Software Changes: Clean or Buggy?

Sunghun KIM, E. James WHITEHEAD, Jr., *Member, IEEE*, and Yi ZHANG, *Member, IEEE*

**Abstract**— This paper introduces a new technique for finding latent software bugs called change classification. Change classification uses a machine learning classifier to determine whether a new software change is more similar to prior buggy changes, or clean changes. In this manner, change classification predicts the existence of bugs in software changes. The classifier is trained using features (in the machine learning sense) extracted from the revision history of a software project, as stored in its software configuration management repository. The trained classifier can classify changes as buggy or clean with 78% accuracy and 65% buggy change recall (on average). Change classification has several desirable qualities: (1) the prediction granularity is small (a change to a single file), (2) predictions do not require semantic information about the source code, (3) the technique works for a broad array of project types and programming languages, and (4) predictions can be made immediately upon completion of a change. Contributions of the paper include a description of the change classification approach, techniques for extracting features from source code and change histories, a characterization of the performance of change classification across 12 open source projects, and evaluation of the predictive power of different groups of features.

**Index Terms**-- Maintenance, Software metrics, Software fault diagnosis, Configuration management, Classification, Association rules, Data mining.

## I. INTRODUCTION

Consider a software developer working on a long duration software product. Like most software developers, she typically makes changes that are *clean*, not containing a latent bug. At times she makes changes that introduce new features, or adapts the software to a changing operational environment. Sometimes she makes *bug-fix* changes that repair a bug. Occasionally she makes a *bug-introducing* change and injects incorrect statements into the source code. Since developers typically do not know they are writing incorrect software, there is always the question of whether the change they just made has introduced a bug.

A bug in the source code leads to an unintended state within the executing software, and this corrupted state eventually results in an undesired external behavior. This is logged in a bug report message in a change tracking system, often many months after the initial injection of the bug into the software. By the time a developer receives the bug report, she must spend time to

Manuscript received March 10, 2006. This work is supported by NSF Grant CCR-01234603, and a Cooperative Agreement with NASA Ames Research Center.

S. Kim is with Massachusetts Institute of Technology, Cambridge, MA USA (telephone: 617-253-1947, e-mail: hunkim@csail.mit.edu).

E. J. Whitehead, Jr. is with University of California, Santa Cruz, CA 95064 USA (telephone: 831-459-1227, e-mail: ejw@cs.ucsc.edu).

Y. Zhang is with University of California, Santa Cruz, CA 95064 USA (telephone: 831-459-4549, e-mail: yiz@soe.ucsc.edu).

reacquaint themselves with the source code and the recent changes made to it. This is time consuming.

If there were a tool that could accurately predict whether a change was buggy or clean immediately after a change was made, it would permit developers to take steps to fix the introduced bugs immediately. Several bug finding techniques could be used, including code inspections, unit testing, and use of static analysis tools. Since these steps would be taken right after a code change was made, the developer would still retain the full mental context of the change. This holds promise for reducing the time required to find software bugs, as well as reducing the time that bugs stay resident in software before removal.

This paper presents a new technique, called *change classification*, for predicting bugs in file-level software changes (a set of line ranges in a single file changed since the previous revision) using machine learning classification algorithms. A key insight behind this work is viewing bug prediction in changes as a kind of classification problem, that is, assigning each change made into one of two classes: clean changes or buggy changes.

The change classification technique involves two steps: training and classification. The change classification algorithms learn from a training set, i.e., a collection of changes that are known to belong to an existing class, that is, the changes are *labeled* with the known class. Features are extracted from the changes, and the classification algorithm learns which features are most useful for discriminating among the various classes. In this context, feature often means some property of the change, such as the frequency of words that are present in source code. This is different from the typical software engineering notion of feature as a software functionality.

Unfortunately for software change classification, we do not have a known standard corpus like the UCI Repository of machine learning [37] or the Reuters Corpus [23] which is commonly

used for evaluation in the text classification domain. Instead, in this paper the file change histories of 12 open source projects are extracted from their software configuration management systems (SCMs), and then features are mined from each change, creating the corpus used to evaluate change classification per project. Each project's features are used to train an Support Vector Machine (SVM) [14] classifier for that project. SVMs are a class of high performance machine learning classification algorithms that have good behavior for a range of text classification problems. After training an SVM using change data from revision  $l$  to  $n$ , if there is a new and unclassified change (i.e., revision  $n+1$ ), this change can be classified as either buggy or clean using the trained classifier model. This act of classification has the effect of predicting which changes have bugs.

This paper makes the following contributions:

*New bug prediction technique: change classification.* Use of machine learning classifiers, as trained on software evolution data and applied to software changes (instead of entire files, functions, or methods), provide a new method for predicting the location of latent software bugs.

*Evaluation of the performance of change classification.* The classification accuracy, recall, and precision are evaluated for each project. An SVM classifies file changes as buggy or clean with 78% accuracy on average (ranging by project from 64%-93% as shown in Figure 2) and 65% buggy change recall on average (43%-98% as shown in Figure 2). This is on par with the best recent work on bug prediction in the research literature [12] [39], with the added benefit that the granularity of prediction is smaller: it is localized to the section of text related to a change instead of a whole file or function/method.

*Techniques for feature extraction from source code and change histories.* In order to extract features from a software evolution history, new techniques for feature extraction were developed.

These have utility for any researcher performing classification experiments using evolution data.

*Evaluation of the performance of individual features, and groups of features.* Since the choice of features can affect the performance of classifiers, each feature's discriminative power for performing change classification is compared. This is performed by evaluating which set of features yields the best overall classification accuracy and recall, and also by examining the relative contributions of individual features.

The remainder of the paper begins with a comparison to related work (Section II). Following is an overview of the approach used to create a corpus, perform change classification, and evaluate its performance (Section III). The process used to create the 12 project corpus is described in detail (Section IV), followed by a brief overview of the SVM algorithm, and the evaluation measures used in this paper. Results from applying change classification to the corpus using all features are presented in Section VI, and Section VII gives an evaluation of the relative importance of the feature groups, and individual features. Section VIII provides discussion of these results, and threats to validity. Section IX concludes the paper.

## II. RELATED WORK

The goal of change classification is to use a machine learning classifier to predict bugs in changes. As a result, related work exists in the area of bug prediction, as well as algorithms for source code clustering and for text classification.

### A. *Predicting Buggy and High Risk Modules*

There is a rich literature for bug detection and prediction. Existing work falls into one of three categories, depending on the goal of the work. The goal of some work is to identify a problematic module list by analyzing software quality metrics or a project's change history [13, 15, 16, 21, 39-41]. This identifies those modules that are most likely to contain latent bugs, but

provides no insight into how many faults may be in each module. Other efforts address this problem, predicting the bug density of each module using its software change history [11, 36]. Work that computes a problematic module list or that determines a fault density are good for determining where to focus quality assurance efforts, but do not provide specific guidance on where, exactly, in the source code to find the latent bugs. In contrast, efforts that detect faults by analyzing source code using static or dynamic analysis techniques can identify specific kinds of bugs in software, though generally with high rates of false positives. Common techniques include type checking, deadlock detection, and pattern recognition [8, 26, 48].

Classification or regression algorithms with features such as complexity metrics, cumulative change count, or bug count are widely used to predict risky entities. Similar to our work, Gyimothy et al. use machine learning algorithms to predict fault classes in software projects [12]. They employ decision trees and neural networks using object-oriented metrics as features to predict fault classes of the Mozilla project across several releases (1.0-1.6). Recall and precision reported in [12] are about 70%, while our change classification accuracy for the Mozilla project is somewhat higher at 77.3% with lower precision at 63.4%. These results reported by Gyimothy et al. are strong. However, they predict faults at the class level of granularity (usually entire files), while the prediction granularity of change classification is much finer, file level *changes*, which for the projects we analyzed average 20 lines of code (LOC) per change. This is significant, since developers need to examine an order of magnitude smaller number of lines of code to find latent bugs with the change classification approach. Gyimothy et al. use release-based classes for prediction, where a release is an accumulation of many versions, whereas change classification applies to the changes between successive individual versions. This allows change classification to be used in an ongoing, daily manner, instead of just for releases, which occur on months-long

time scales.

Kim et al. proposed BugMem to capture fault patterns in previous fixes and to predict future faults using captured fix memories [17]. Mizuno and Kikuno use an email spam filter approach to capture patterns in fault-prone software modules [31]. These two approaches are similar to change classification in that they learn fault patterns and predict future faults based on them. However, they classify static code (such as the current version of a file) while our approach classifies file changes. The patterns they can learn are limited to source code only, whereas change classification uses features from all possible sources such as source code, metadata, delta, complexity metrics, and change logs.

Brun and Ernst [5] use two classification algorithms to find hidden code errors. Using Ernst's Daikon dynamic invariant detector, invariant features are extracted from code with known errors and with errors removed. They train SVM and decision tree classifiers using the extracted features, then classify invariants in the source code as either fault-invariant or non-fault-invariant. The fault-invariant information is used to find hidden errors in the source code. Reported classification accuracy is 10.6% on average (9% for C and 12.2% for Java), with classification precision of 21.6% on average (10% for C and 33.2% for Java), and the best classification precision (with top 80 relevant invariants) of 52% on average (45% for C and 59% for Java). The classified fault invariants guide developers to find hidden errors. Brun and Ernst's approach is similar to our work in that they try to capture properties of buggy code and use it to train machine learning classifiers to make future predictions. However, they used only invariant information as features, which leads to lower accuracy and precision. In contrast, change classification uses a broader set of features including source code, complexity metrics, and change metadata.

Hassan and Holt use a caching algorithm to compute the set of fault prone modules, called the top-ten list [13]. They use four factors to determine this list: software that was most frequently modified, most recently modified, most frequently fixed, and most recently fixed. Kim et al. proposed the bug cache algorithm to predict future faults based on previous fault localities [18]. Ostrand et al. identified the top 20% of problematic files in a project [39, 40]. Using future fault predictors and a negative binomial linear regression model, they predict the fault density of each file.

Khoshgoftaar and Allen have proposed a model to list modules according to software quality factors such as future fault density [15, 16, 21]. The inputs to the model are software complexity metrics such as LOC, number of unique operators, and cyclomatic complexity. A step-wise multi regression is then performed to find weights for each factor [15, 16]. Mockus and Weiss predict risky modules in software using a regression algorithm and change measures, such as number of systems touched, number of modules touched, number of lines of added code, and number of modification requests [33].

Pan et al. use metrics computed over software slice data in conjunction with machine learning algorithms to find bug-prone software files or functions [41]. Their approach tries to find faults in the whole code, while our approach focuses on file changes.

### *B. Mining Buggy Patterns*

One thread of research attempts to find buggy or clean code patterns in the history of development of a software project.

Williams and Hollingsworth use project histories to improve existing bug finding tools [51]. Using a return value without first checking its validity may be a latent bug. In practice, this approach leads to many false positives, as typical code has many locations where return values

are used without checks. To remove the false positives, Williams and Hollingsworth use project histories to determine which kinds of function return values must be checked. For example, if the return value of *foo* was always verified in the previous project history, but was not verified in the current source code, it is very suspicious. Livshits and Zimmermann combined software repository mining and dynamic analysis to discover common use patterns and code patterns that are likely errors in Java applications [25]. Similarly, PR-Miner mines common call sequences from a code snapshot, and then marks all non-common call patterns as potential bugs [24].

These approaches are similar to change classification, since they use project specific patterns to determine latent software bugs. However, the mining is limited to specific patterns such as return types or call sequences, and hence limits the type of latent bugs that can be identified.

### *C. Classification, Clustering, Associating, and Traceability Recovery*

Several research efforts share a similarity with bug classification in that they also extract features (terms) from source code, and then feed them into classification or clustering algorithms. These efforts have goals other than predicting bugs, including classifying software into broad functional categories [19], clustering related software project documents [20, 28], and associating source code to other artifacts such as design documents [29].

Krovtez et al. use terms in the source code (as features) and SVM to classify software projects into broad functional categories such as communications, databases, games, and math [19]. Their insight is that software projects in the same category will share terms in their source code, thereby permitting classification.

Research that categorizes or associates source code with other documents (traceability recovery) is similar to ours in that it gathers terms from the source code and then uses learning or statistical approaches to find associated documents [2, 42]. For example, Maletic et al. [28, 29]



extracted all features available in the source code via Latent Semantic Analysis (LSA), then used this data to cluster software and create relationships between source code and other related software project documents. In a similar vein, Kuhn et al. used partial terms from source code to cluster the code to detect abnormal module structures [20]. Antoniol et al. used stochastic modeling and Bayesian classification for traceability recovery [2]. Their work differs from our work in that they only use features from source code, while our change classification learns from project history data, including change deltas, change log text, and authors. Traceability recovery focuses on finding associations among source code and other documents, while change classification tries to identify each change as buggy or clean.

Similar in spirit to change classification is work that classifies bug reports or software maintenance requests [3, 10]. In their research [3, 10], keywords in bug reports or change requests are extracted and used as features train a machine learning classifier. The goal of the classification is to place a bug report into a specific category of bug report, or to find the developer best suited to fix a bug. This work, along with change classification, highlights the potential of using machine learning techniques in software engineering. If an existing concern – such as assigning bugs to developers, can be recast as a classification problem, then it is possible to leverage the large collection of data stored in bug tracking and software configuration management systems.

#### *D. Text Classification*

Text classification is a well-studied area with a long research history. Using text terms as features, researchers have proposed many algorithms to classify text documents [46], such as classifying news articles into their corresponding genres. Among existing work on text classification, spam filtering [16] is the most similar to ours. Spam filtering is a binary

classification problem to identify email as spam or ham (not spam). This paper adapts existing text classification algorithms into the domain of source code change classification. Our research focuses on generating and selecting features related to buggy source code changes.

#### *E. Summary*

Change classification differs from previous bug prediction work since it:

**Classifies changes:** Previous bug prediction work focuses on finding prediction or regression models to identify fault-prone or buggy modules, files, and functions [11, 36, 38]. Change classification predicts whether there is a bug in any of the lines that were changed in one file in one SCM commit transaction. This can be contrasted with making bug predictions at the module, file, or method level. Bug predictions are immediate, since change classification can predict buggy changes as soon as a change is made.

**Uses bug-introducing changes:** Most bug prediction research uses bug-fix data when making predictions or validating their prediction model. Change classification uses bug-introducing changes, which contains the exact commit/line changes that injected a bug, who introduced it, and the time it occurred. Bug-fix changes indicate only roughly where the bug occurred. Bug-introducing changes allow us to label changes as buggy or clean with information about the source code at the moment a bug was introduced.

**Uses features from source code:** When selecting predictors, bug prediction research usually does not take advantage of the information directly provided by the source code, and thereby miss a valuable source of features. Change classification uses every term in the source code—every variable, method call, operator, constant, comment text, and more—as features to train our change classification models.

**Is independent of programming language:** Our change classification approach is

programming language independent, since we use a bag-of-words method [45] for generating features from source code. The projects we analyzed span many popular current programming languages, including C/C++, Java, Perl, Python, Java Script, PHP, and XML.

### III. OVERVIEW OF CHANGE CLASSIFICATION APPROACH

The primary steps involved in performing change classification on a single project are outlined below.

#### **Creating a Corpus:**

1. File level changes are extracted from the revision history of a project, as stored in its software configuration management (SCM) repository. (Described further in Section IV.A).
2. The bug fix changes for each file are identified by examining keywords in SCM change log messages, part of the data extracted from the SCM repository in step 1 (Section IV.B).
3. The bug-introducing and clean changes at the file level are identified by tracing backwards in the revision history from bug fix changes, using SCM annotation information (Section IV.B).
4. Features are extracted from all changes, both buggy and clean. Features include all terms in the complete source code, the lines modified in each change (delta), and change metadata such as author and change time. Complexity metrics, if available, are computed at this step. Details on the feature extraction techniques used are presented in Section IV.C.

At the end of step 4 a project-specific corpus has been created, a set of labeled changes with a set of features associated with each change [52].

#### **Classification:**

5. Using the corpus, a classification model is trained. While many classification techniques could be employed, this paper focuses on the use of SVM, outlined in Section V.A.

6. Once a classifier has been trained, it is ready to use. New changes can now be fed to the classifier, which determines whether a new change is more similar to a buggy change, or a clean change.

Machine learning classifiers have varying performance depending on the characteristics of the data set used to train the classifier, and the information available in the text being classified. This paper examines the behavior of change classification by assessing its predictive performance on 12 open source software systems. Since inclusion and omission of different feature sets can affect the predictive performance, an examination of the performance of different feature groups is also performed. The overall approach for these two evaluations is provided below:

### **Evaluation of Change Classification**

1. Classification performance is evaluated using the 10-fold cross-validation method [35] and computation of the standard classification evaluation measures of accuracy, recall, precision, and F-value. Definitions of these measures are provided in Sections V.B and V.C, and the actual measured values are presented in Section VI.A.
2. Since there are only two outcomes, a prediction approach that randomly guesses whether a change is buggy or clean (a dummy classifier) might have better performance than use of a machine learning classifier. Recall-precision curves are presented in Section VI.B, along with analysis of the performance of change classification relative to the dummy classifier.

### **Evaluation of Feature Group Performance:**

1. Subsets of the corpus are created by combining features, then retraining the classifier (change classification step 5 above) and recomputing the classification performance (evaluation of change classification step 1 above). This permits a comparison of classification performance using different feature combinations (Sections VII.A).
2. Using the chi-square measure which is commonly used for feature selection [54], each

feature is ranked to determine how correlated each feature is to the buggy and clean classes.

This provides insight into which features are the most informative for change classification (Section VII.B).

#### IV. CREATING THE CORPUS

##### A. Change History Extraction

Kenyon [4] is a system that automates the extraction of source code change histories from SCM systems such as CVS and Subversion. Kenyon automatically checks out a user-defined range of revisions stored in a project's SCM repository, reconstructing logical transactions from individual file commits for the CVS system [55]. Revisions include files and their change information. From checked out revisions, Kenyon extracts change information such as the change log, author, change date, source code, change delta, and change metadata. This information is then fed into the feature extraction process to convert a file change into a vector of features.

**Table 1. Subject programs and Summary of corpus information.**

Project	Revisions.	Period	# of clean changes	# of buggy changes	% of buggy changes	# of features
Apache HTTP 1.3 (A1)	500-1000	10/1996-01/1997	579	121	17.3	11,445
Bugzilla (BUG)	500-1000	03/2000-08/2001	149	417	73.7	10,148
Columba (COL)	500-1000	05/2003-09/2003	1,270	530	29.4	17,411
Gaim(GAI)	500-1000	08/2000-03/2001	742	451	37.8	9,281
GForge(GFO)	500-1000	01/2003-03/2004	339	334	49.6	8,996
Jedit (JED)	500-750	08/2002-03/2003	626	377	37.5	13,879
Mozilla (MOZ)	500-1000	08/2003-08/2004	395	169	29.9	13,648
Eclipse(ECL)	500-750	10/2001-11/2001	592	67	10.1	16,192
Plone(PLO)	500-1000	07/2002-02/2003	457	112	19.6	6,127
PostgreSQL (POS)	500-1000	11/1996-02/1997	853	273	24.2	23,247
Scarab (SCA)	500-1000	06/2001-08/2001	358	366	50.5	5,710
Subversion (SVN)	500-1000	01/2002-03/2002	1,925	288	13.0	14,856
Total	N/A	N/A	8,285	3,505	29.7	150,940

One challenge in change classification is ensuring the memory used by the SVM classifier does not grow too large. To limit the memory footprint, only a subset of each project's revisions were selected, typically 500 revisions. To try to ensure projects were compared at more-or-less

the same level of maturity, file change features were extracted from revisions 500-1000 (or revisions 500-750 for big projects). Additionally, there was some initial concern that the change patterns in the first part of a project (revisions 1-500) may not be stable ([13] noted similar concerns), but later analysis showed this was not the case (in general, project maturity level has no substantive impact on prediction results). Table 1 provides an overview of the projects examined in this research, the range of revisions extracted, and the real-world duration of each range.

**Table 2. Average LOC.**

Project	File	File change	Function/ Method	Function/Method change
Apache HTTP 1.3 (A1)	455.73	15.42	28.32	15.83
Bugzilla (BU)	375.37	18.30	N/A	N/A
Columba (CO)	143.3	14.94	15.64	10.99
Gaim(GA)	832	19.64	38.43	11.09
GForge(GFO)	155.49	17.73	N/A	N/A
JEdit (JED)	325.78	23.64	18.74	7.65
Mozilla (MOZ)	285	21.20	N/A	N/A
Eclipse(ECL)	230.29	48.26	16.9	13.77
Plone(PLO)	49.11	9.7	N/A	N/A
PostgreSQL (POS)	282.92	14.28	32.21	25
Scarab (SCA)	145.98	21.75	16.06	15.21
Subversion (SVN)	354.31	15.35	33.81	11.923
Average	302.94	20.02	25.01	13.93

One of the advantages of classifying file changes was discussed earlier: it provides predictions at a fine level of granularity (a single change to a single file). For the 12 projects examined in this paper, Table 2 shows average values for the number of LOC in a file change, in each file, in a function/method change, and in each function/method. On average, the number of LOC per file change is 20 while the average LOC per file is 300. For example, if a tool predicts bugs at the file level, it is necessary to inspect 300 LOC on average to locate the line(s) containing the bug. Since our approach classifies file changes, the prediction is at the file change level, and hence only 20 lines on average need to be inspected.

### B. Identifying Bug-introducing Changes

The first step toward identifying bug-introducing changes is to find bug-fix changes. Bug-fix changes are identified by mining change log messages. Two approaches are used for this step: searching for keywords such as "Fixed" or "Bug" [32] and searching for references to bug reports like "#42233" [6, 7, 47]. If a keyword or bug report reference is found, the changes in the associated commit comprise a bug-fix. Table 3 lists keywords or phrases used to identify bug-fix commits. Manual verification of identified bug-fix changes is recommended to ensure the selected keywords or phrases are correctly identifying bug-fix changes. For example, if a commit log stated, "This is not a bug fix", its commit should not be identified as a fix. For the systems studied in this paper, one of the authors manually verified that the identified fix commits were, indeed, fixes.

**Table 3. Keywords and reference identifiers to find fix commits. \* bug id reference is a 7-digit number.**

Project	Keywords or phrases
Apache HTTP 1.3 (A1)	Patch, fix, bug
Bugzilla (BUG)	Fix, bug, * bug id reference number
Columba (COL)	[bug], [bugfix]
Gaim (GAI)	Patch, fix, bug
GForge (GFO)	Patch, fix, bug
Jedit (JED)	Patch, fix, bug
Mozilla (MOZ)	* bug id reference number
Eclipse JDT (ECL)	* bug id reference number
Plone (PLO)	Patch, fix, bug
PostgreSQL (POS)	Patch, fix, bug
Scarab (SCA)	Patch, fix, bug, issue number
Subversion (SVN)	Fixed issue number

One potential issue for identifying bug-fixes using bug tracking system identifiers is the common use of bug tracking systems to record both bug reports and new feature additions. This causes new feature changes to be identified as bug-fix changes. Among the systems studied in this paper, Bugzilla and Scarab both use bug tracking systems to record new feature additions, and as result the percentage of buggy changes found in these systems is higher than for other

projects (Bugzilla, 73.7%, and Scarab, 50.5%). For these two systems, what this paper terms a “buggy” change should be interpreted as a “buggy or new feature” change. Similarly, for these two systems, predictions of buggy changes should be interpreted as predictions of “buggy or new feature” changes.

Once a commit has been determined to contain a fix, it is possible to work backwards in the revision history to determine the initial bug introducing change. The bug-introducing change identification algorithm proposed by Śliwerski, Zimmermann, and Zeller (SZZ algorithm) is used in this paper [47]. After identifying bug fixes, SZZ uses a *diff* tool to determine what changed in the bug-fixes. The diff tool returns a list of regions that differ in the two files; each region is called a “hunk”. It observes each hunk in the bug-fix and assumes that the deleted or modified source code in each hunk is the location of a bug.

Revision 1 (by kim, bug-introducing)		Revision 2 (by ejw)	
1 kim	1: public void bar() {	2 ejw	1: public void foo() {
1 kim	2: // print report	1 kim	2: // print report
1 kim	3: if (report == null) {	1 kim	3: if (report == null){
1 kim	4: println(report);	2 ejw	4: println(report.str);
1 kim	5: }	1 kim	5: }

Revision 3 (by kai, bug-fix)	
2 ejw	1: public void foo() {
1 kim	2: // print report
3 kai	3: if (report != null) {
1 kim	4: println(report);
1 kim	5: }

**Figure 1. Example bug-fix and source code changes.** A null-value checking bug is injected in revision 1, and fixed in revision 3.

Finally, SZZ tracks down the origins of the deleted or modified source code in the hunks using the built-in *annotate* functionality of SCM systems. The annotate feature computes, for each line in the source code, the most recent revision in which the line was changed, and the developer who made the change. The discovered origins are identified as bug-introducing changes.

Figure 1 shows an example of the history of development of a single function over three



revisions:

- Revision 1 shows the initial creation of function *bar*, and the injection of a bug into the software, the line `if (report == null) {` which should be `!=` instead. The leftmost column of each revision shows the output of the SCM annotate command, identifying the most recent revision for each line and the developer who made the revision. Since this is the first revision, all lines were first modified at revision 1 by the initial developer *'kim.'* The second column of numbers in revision 1 lists line numbers within that revision.
- In the second revision, two changes were made. The function *bar* was renamed to *foo*, and `println` has argument *'report.str'* instead of *'report.'* As a result, the annotate output shows lines 1 and 4 as having been most recently modified in revision 2 by *'ejw.'*
- Revision 3 shows a change, the actual bug-fix, changing line 3 from `==` to `!=`.

The SZZ algorithm then identifies the bug-introducing change associated with the bug-fix in revision 3. It starts by computing the delta between revisions 3 and 2, yielding line 3. SZZ then uses SCM annotate data to determine the initial origin of line 3 at revision 2. This is revision 1, the bug-introducing change.

One assumption of the presentation so far is that a bug is repaired in a single bug-fix change; what happens when a bug is repaired across multiple commits? There are two cases. In the first case, a bug repair is split across multiple commits, with each commit modifying a separate section of code (code sections are disjoint). Each separate change is tracked back to its initial bug-introducing change, which is then used to train the SVM classifier. In the second case, a bug fix occurs incrementally over multiple commits, with some later fixes modifying earlier ones (fix code partially overlaps). The first patch in an overlapping code section would be traced back to the original bug-introducing change. Later modifications would not be traced back to the original

bug-introducing change; instead they would be traced back to an intermediate modification, which is identified as bug-introducing. This is appropriate, since the intermediate modification did not correctly fix the bug, and hence is simultaneously a bug-fix, and buggy. In this case, the classifier is being trained with the attributes of the buggy intermediate commit, a valid bug-introducing change.

### *C. Feature Extraction*

To classify software changes using machine learning algorithms, we need to train a classification model must be trained using features of buggy and clean changes. In this section, we discuss techniques for extracting features from a software project change history.

A file change involves two source code revisions (an old revision and a new revision) and a change delta that records the added code (added delta) and deleted code (deleted delta) between the two revisions. A file change has associated metadata, including the change log, author, and commit date. By mining change histories, we can derive features such as co-change counts to indicate how many files are changed together in a commit, the number of authors of a file, and previous change count of a file. Every term in the source code, change delta, and change log texts are used as features. We detail our feature extraction method below.

#### *1) Feature Extraction from Change Metadata*

We gather 8 features from change metadata: author, commit hour (0, 1, 2, ... 23), commit day (Sunday, Monday, ..., Saturday), cumulative change count, cumulative bug count, length of change log, changed LOC (added delta LOC + deleted delta LOC), and new revision source code LOC. In other research, cumulative bug and change counts are commonly used as bug predictors [11, 33, 38, 40, 47, 48, 56].

## 2) *Complexity Metrics as Features*

Software complexity metrics are commonly used to measure software quality and predict defects in software modules [12, 15, 30]. Modules with higher complexity measures tend to correlate with greater fault incidence. We compute a range of traditional complexity metrics of source code using the Understand C/C++ and Java tools [44]. As a result, we extract 61 complexity metrics (every complexity metric these tools compute) for each file including LOC, lines of comments, cyclomatic complexity, and max nesting. Since we have two source code files involved in each change (old and new revision files), we compute and use as features the difference in value, a complexity metric delta for each complexity metric between these two revisions.

## 3) *Feature Extraction from Change Log Messages, Source Code, and File names*

Change log messages are similar to email or news articles in that they are human readable texts. Each word in a change log message carries meaning. Feature engineering from texts is a well studied area, with the bag-of-words, latent semantic analysis (LSA), and vector models being widely used approaches for text classification [43, 45]. Among them, the bag-of-words (BOW) approach, which converts a stream of characters (the text) into a bag of words (index terms), is simple and performs fairly well in practice [45, 46]. We use BOW to generate features from change log messages.

We extract all words except for special characters, and convert all words to lowercase. The existence (binary) of a word in a document is used as a feature. Although stemming (removing stems) and stopping (removing very frequent words) are used by researchers in the text classification community to reduce the number of features, we did not perform these steps to simplify our experiments. Additionally, use of stemming on variable, method, or function names

is generally inappropriate, since this changes the name.

We use every term in the source code as features, including operators, numbers, keywords, and comments. To generate features from source code, we use a modified version of BOW, called BOW+, that extracts operators in addition to all terms extracted by BOW, since we believe operators such as “!=”, “++”, and “&&” are important terms in source code. We perform BOW+ extraction on added delta, deleted delta, and new revision source code. This means that every variable, method name, function name, keyword, comment word and operator—everything in the source code separated by whitespace or a semicolon—is used as a feature.

We also convert the directory and file name into features, since they encode both module information and some behavioral semantics of the source code. For example, the file (from the Columba project), *‘ReceiveOptionPanel.java’* in the directory, *‘src/mail/core/org/columba/mail/gui/config/account/’* reveals that the file receives some options using a panel interface, and the directory name shows the source code is related to *‘account’*, *‘configure’*, and *‘graphical user interface’*. Some researchers perform bug predictions at the module granularity by assuming that bug occurrences in files in the same module are correlated [11, 13, 48].

We use the BOW approach by removing all special characters such as slashes, then extracting words in the directory and file names. Directory and file names often use Camelcase, concatenating words then identifying word breaks with capitals [50]. For example, *‘ReceiveOptionPanel.java’* combines *‘receive’*, *‘option’*, and *‘panel’*. To extract such words correctly, we use a case change in a directory or file name as a word separator. We call this method BOW++. Table 4 summarizes features generated and used in this paper.

**Table 4. Feature groups.** Feature group description, extraction method, and example features.

Feature Group	Description	Extraction method	Example Features
Added Delta (A)	Terms in the added delta source code	BOW+	if, while, for, ==
Deleted Delta (D)	Terms in the deleted delta source code	BOW+	true, 0, <, ++, int
Directory/File Name (F)	Terms in the directory/file names	BOW++	src, module, java
Change Log (L)	Terms in the change log	BOW	fix, added, new
New Revision Source Code (N)	Terms in the new revision source code file	BOW+	if, ll, !=, do, while, string, false
Metadata (M)	Change metadata such as time and author	Direct	author: hunkim, commit hour: 12
Complexity Metrics (C)	Software complexity metrics of each source code	Understand tools [44]	LOC: 34, Cyclomatic: 10

#### 4) Feature Extraction Summary

Using the feature engineering technique described previously, features are generated from all file changes in the analyzed range of revisions. Each file change is represented as an instance, a set of features. Using the bug-introducing change identification algorithm, we label each instance as clean or buggy. Table 1 summarizes the corpus information. Consider the Apache 1.3 HTTP server project. For this project, the corpus includes changes in revisions 500-1000, a total of 700 changes, of which 579 are clean and 121 buggy. From the 700 changes, 11,445 features were extracted.

## V. SUPPORT VECTOR MACHINES AND EVALUATION TECHNIQUES

Among many classification algorithms, Support Vector Machine (SVM) [14] is used to implement and evaluate the change classification approach for bug prediction because it is a high performance algorithm that is currently used across a wide range of text classification applications. Several good quality implementations of SVM are readily available; the Weka Toolkit [52] implementation is used in this study. Below, we provide an overview description of SVM, and then describe the measures used in our evaluation of SVM for change classification. There is a substantial literature on SVM; the interested reader is encouraged to pursue [14] or [49] for an in-depth description.

### A. Overview of Support Vector Machines

SVMs were originally designed for binary classification, where the class label can take only two different values. An SVM is a discriminative model that directly models the decision boundary between classes. An SVM tries to find the maximum margin hyperplane, a linear decision boundary with the maximum margin between it and the training examples in class 1 and training examples in class 2 [49]. This hyperplane gives the greatest separation between the two classes.

### B. 10-Fold Cross Validation

Among the labeled instances in a corpus, it is necessary to decide which subset is used as a training set or a test set, since this affects classification accuracy. The 10-fold cross-validation technique [35, 52] is used to handle this problem in our experiment.

### C. Measuring Accuracy, Precision, Recall, and F-Value

There are four possible outcomes from using a classifier on a single change: classifying a buggy change as buggy ( $b \rightarrow b$ ), classifying a buggy change as clean ( $b \rightarrow c$ ), classifying a clean change as clean ( $c \rightarrow c$ ), and classifying a clean change as buggy ( $c \rightarrow b$ ). With a known good set of data (the test set fold that was pulled aside and not used for training), it is then possible to compute the total number of buggy changes correctly classified as buggy ( $n_{b \rightarrow b}$ ), buggy changes incorrectly classified as clean ( $n_{b \rightarrow c}$ ), clean changes correctly classified as clean ( $n_{c \rightarrow c}$ ), and clean changes incorrectly classified as buggy ( $n_{c \rightarrow b}$ ).

Note that the known good dataset is derived by tracing bug fix changes back to bug-introducing changes. The set of bug-introducing (buggy) changes represents those bugs in the code that had sufficiently observable impacts to warrant their repair. The set of bug-introducing changes is presumably smaller than the set of all changes that introduce a bug into the code. The

comprehensive set of all bugs injected into the code during its development lifetime is unknown for the projects examined in this paper. It would require substantial time and effort by a large team of experienced software engineers to develop a comprehensive approximation of the total set of real bugs.

Accuracy, recall, precision, and F value measures are widely used to evaluate classification results [46, 53]. These measures are used to evaluate our file change classifiers, as follows [1, 34,

$$53]: \text{Accuracy} = \frac{n_{b \rightarrow b} + n_{c \rightarrow c}}{n_{b \rightarrow b} + n_{b \rightarrow c} + n_{c \rightarrow c} + n_{c \rightarrow b}}$$

That is, the number of correctly classified changes over the total number of changes. This is a good overall measure of the predictive performance of change classification. Since there are typically more clean changes than buggy changes, this measure could potentially yield a high value if clean changes are being predicted better than buggy changes. Precision and recall measures provide insight into this.

$$\text{Buggy change precision, } P(b) = \frac{n_{b \rightarrow b}}{n_{b \rightarrow b} + n_{c \rightarrow b}}$$

This represents the number of correct classifications of the type ( $n_{b \rightarrow b}$ ) over the total number of classifications that resulted in a bug outcome. Or, put another way, if the change classifier predicts a change is buggy, what fraction of these changes really contains a bug?

$$\text{Buggy change recall, } R(b) = \frac{n_{b \rightarrow b}}{n_{b \rightarrow b} + n_{b \rightarrow c}}$$

This represents the number of correct classifications of the type ( $n_{b \rightarrow b}$ ) over the total number of changes that were actually bugs. That is, of all the changes that are buggy, what fraction does the change classifier predict?

$$\text{Buggy change F1-value} = \frac{2 * P(b) * R(b)}{P(b) + R(b)}$$

This is a composite measure of buggy change precision and recall.

Similarly, clean change recall, precision, and F-value can be computed:

$$\text{Clean change precision, } P(c) = \frac{n_{c \rightarrow c}}{n_{c \rightarrow c} + n_{b \rightarrow c}}$$

If the change classifier predicts a change is clean, what fraction of these changes really is clean?

$$\text{Clean change recall, } R(c) = \frac{n_{c \rightarrow c}}{n_{c \rightarrow c} + n_{c \rightarrow b}}$$

Of all the changes that are clean, what fraction does the change classifier predict?

$$\text{Clean change F1-value} = \frac{2 * P(c) * R(c)}{P(c) + R(c)}$$

This is a composite measure of clean change precision and recall.


## VI. EVALUATION OF CHANGE CLASSIFICATION

This section evaluates change classification in two ways. The first section presents the typical machine learning classifier assessment metrics of accuracy, recall, precision, and F values. These results were computed using the complete set of features extracted for each project. The second section explores whether change classification performs better than just randomly guessing.

### A. Accuracy, Precision, and Recall

Figure 2 shows accuracy, buggy change recall, and buggy change precision of the 12 projects using all features listed in Table 1.





**Figure 2. Change classification accuracy, buggy change recall, and buggy change precision of 12 projects using SVM and all features.**

Detailed recall, precision, and F1 values are reported in Table 5. Buggy change recall ranges between 43% and 86%, and buggy change precision ranges between 44% and 85%. The change classification approach can predict bugs with 64% to 92% accuracy at the file change level of granularity. With a file-level change having on average, 20 LOC, this is the most specific prediction granularity in the literature. Overall, the combined prediction accuracy and granularity of change classification exceed the state of the art reported in the bug prediction literature.

**Table 5. Change classification accuracy, recall, precision, and F1 values of 12 open source projects. SVM classification algorithm with all features are used.**


Project	Accuracy	Buggy change recall	Buggy change precision	Buggy change F1	Clean change recall	Clean change precision	Clean change F1
A1	0.85	0.5	0.56	0.53	0.92	0.9	0.91
BUG	0.78	0.86	0.85	0.86	0.56	0.6	0.58
COL	0.76	0.58	0.59	0.59	0.83	0.83	0.83
GAI	0.71	0.63	0.62	0.62	0.76	0.77	0.77
GFO	0.64	0.6	0.60	0.6	0.67	0.67	0.67
JED	0.65	0.53	0.54	0.53	0.73	0.72	0.72
MOZ	0.77	0.57	0.63	0.6	0.86	0.83	0.84
ECL	0.92	0.61	0.61	0.61	0.96	0.96	0.96
PLO	0.8	0.48	0.49	0.48	0.89	0.87	0.87
POS	0.73	0.43	0.44	0.43	0.82	0.82	0.82
SCA	0.79	0.78	0.8	0.79	0.8	0.78	0.79
SVN	0.9	0.59	0.6	0.6	0.94	0.94	0.94

### B. Comparing SVM with a Dummy Predictor

There are tradeoffs between precision and recall, and it is often possible to improve recall by reducing precision and vice versa. Most machine learning classifiers use a threshold value to

classify instances. For example, SVM uses the distance between each instance and the hyperplane to measure the weights of each instance. If an instance's weight is greater than the threshold value, the instance belongs to class 1, otherwise it belongs to class 2. By lowering or raising the threshold, it is possible to change recall and precision. Usually by lowering recall, precision can be increased. For example, buggy change recall can easily go up to 100% by predicting all changes as buggy, but the precision will be very low.

A recall-precision curve shows the trade-off between recall and precision. Figure 3 gives the SVM classifier recall-precision curves of three selected projects, Bugzilla, Mozilla, and Scarab (in solid lines). The curve for the Bugzilla project shows that the precision grows up to about 95% (with 20% recall). For Mozilla and Scarab, the precision can reach 85-90% by lowering the recall to 20%.



**Figure 3.** Buggy change recall-precision curves of selected 3 projects, Bugzilla, Mozilla, and Scarab using SVM. Dummy classifier recall-precision curves are shown as dotted lines while the solid lines represent SVM classifier recall-precision curves.

How is SVM recall-precision better than other approaches, such as randomly guessing changes (a dummy classifier) as buggy or clean? Since there are only two classes, the dummy classifier may work well. For example, 73.7% of Bugzilla changes are buggy. By predicting all changes as buggy, buggy recall would be 100% and precision would be 73.7%. Is this better than the results when using SVM? The recall-precision curves of the dummy (dotted lines) and SVM (solid lines) classifiers of three selected projects are compared in Figure 3. Precision for the Bugzilla dummy classifier is stuck at 73.7%, while SVM precision grows up to 95% (with 20% recall). Similarly for other projects, SVM can improve buggy change precision by 20% - 35%.

### *C. Correlation between Percentage of Bug Introducing Changes and Classification Accuracy*

One observation that can be made from Table 1 is that the percentage of changes that are buggy varies substantially among projects, ranging from 10.1% of changes for Eclipse to 73.7% for Bugzilla. One explanation for this variance is the varying use of change log messages among projects. Bugzilla and Scarab, being change tracking tool projects, have a higher overall use of change tracking. It is likely that for those projects, the class of buggy changes also encompasses other kinds of modifications. For these projects, change classification can be viewed as successfully predicting the kinds of changes that result in change tracking tool entries.

**Table 6. Correlation between the percentages of buggy changes and change classification performance.**

	Buggy % vs. accuracy	Buggy % vs. buggy change recall	Buggy % vs. buggy change precision
Correlation	-0.56	0.77	0.64

One question that arises is whether the percentage of buggy changes for a project affects change classification performance such as accuracy, recall, and precision? A Pearson correlation was computed between the percentage of buggy changes, and the measures of accuracy, recall, and precision for the 12 projects analyzed in this paper. Table 6 lists the correlation values. A

correlation value of 1 indicates tight correlation, while 0 indicates no correlation. The values show a weak negative correlation for accuracy, and weak, but not significant correlations for buggy recall and precision.

## VII. EVALUATION OF FEATURES AND FEATURE GROUPS

### A. Change Classification Using Selected Feature Groups

This section evaluates the accuracy of different feature group combinations for performing change classification. First, a classification model is trained using features from one feature group, and then its accuracy and recall are measured. Following, a classification model is trained using all feature groups except the one feature group, with accuracy and recall measured for this case as well. In addition, an evaluation is made focusing on the combination of features extracted solely from the source code (added delta, new revision source code, and deleted delta).

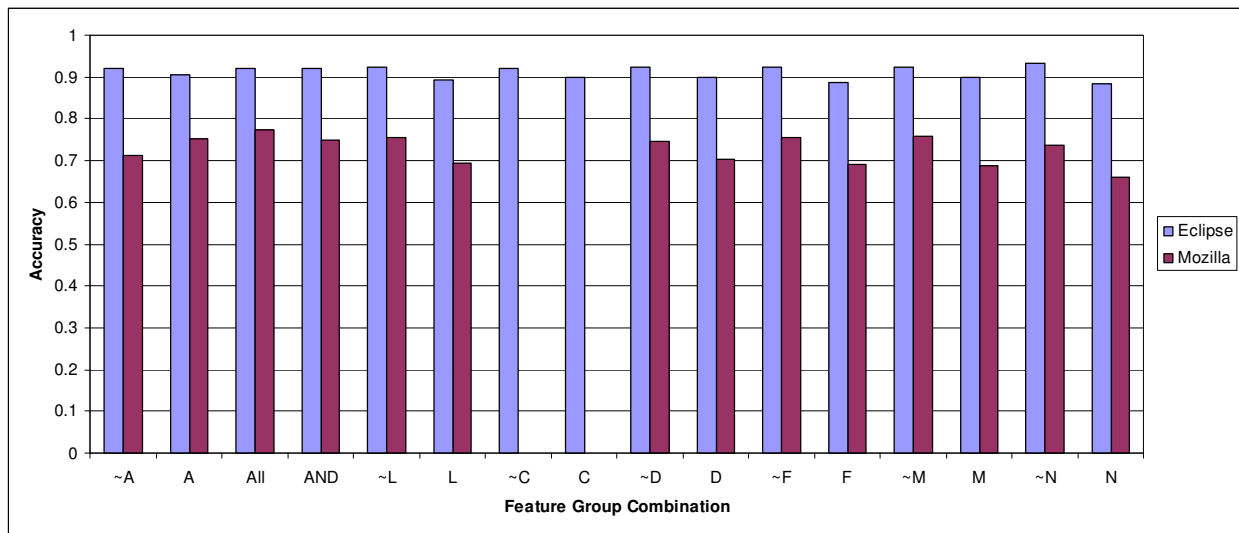
**Table 7. Feature groups.** The number of features in each feature group is shown for all analyzed projects.

Feature Group	Number of features of projects											
	A1	BUG	COL	GAI	GFO	JED	MOZ	ECL	PLO	POS	SCA	SVN
Added Delta (A)	2024	2506	3811	2094	1895	2939	3079	2558	1540	3532	1290	2663
Deleted Delta (D)	1610	1839	3227	1956	1832	2352	2176	2200	1073	2995	836	2117
Directory/File Name (F)	93	66	559	39	242	377	105	456	221	472	106	195
Change Log (L)	1257	1124	869	1094	3970	431	959	53	2835	1161	650	2474
New Revision Source Code (N)	6330	4604	8814	3967	4481	7649	7320	10794	2671	14956	2697	7276
Metadata (M)	8	8	8	8	8	8	8	8	8	8	8	8
Complexity Metrics (C)	122	0	122	122	0	122	0	122	0	122	0	122
*Total	11445	10148	17411	9281	8996	13879	13648	16192	6127	23247	5710	14856

Extracted features are organized into groups based on their source. For example, features extracted just from the change log messages are part of the Change Log (L) feature group. Table 7 provides a summary of feature groups and the number of features in each group. Software complexity metrics were computed only for C/C++ and Java source code, since tools were not available for computing these metrics for Java Script, Perl, PHP, and Python.


Figure 4 shows the change classification accuracy for the Mozilla and Eclipse projects using various feature group combinations. The abbreviations for each feature group are shown in Table

7. An abbreviation means only the feature group is used for classification. For example, ‘D’ means only features from the Deleted Delta group were used. The ‘~’ mark indicates the corresponding feature group is excluded. For example, ‘~D’ means all features were used except for D (Deleted delta). The feature group “AND” is the combination of all source code feature groups (A, N, and D). The accuracy trend of the two projects is different, but they share some properties. For example, the accuracy using only one feature group is lower than using multiple feature groups.



**Figure 4. Feature group combination accuracy for Eclipse and Mozilla using SVM. Complexity metrics (C) are not available for Mozilla, so ~C and C are omitted.**

The average accuracy of 12 open source projects using various feature combinations is shown in Figure 5. Using a feature combination of only source code (A, N, and D combined) leads to a relatively high accuracy, while using only one feature group from the source code, such as A, N, or D individually, does not lead to high accuracy. Using only ‘L’ (change log features) leads to the worst accuracy.



**Figure 5. Average feature group combination accuracy across the 12 analyzed projects using SVM**

After analyzing the combinations of feature groups, the feature combination that yields the best accuracy and best recall for each project are identified, as shown in Table 8. The results indicate that there is no feature combination that works best across projects, and that frequently the feature group providing the best accuracy is not the same as the feature group providing the best buggy recall. A practical implication of this data is that each project has the opportunity to engage in a project-specific feature selection process that can optimize for either accuracy or recall, but often not both simultaneously.

**Table 8. Feature group combination yielding the best classification accuracy and buggy change recall using SVM.** The feature group listed in parentheses yields the best accuracy/buggy change recall.

Projects	A1	BUG	COL	GAI	GFO	JED	MOZ	ECL	PLO	POS	SCA	SVN
Best accuracy	0.86 (~F)	0.79 (~M)	0.77 (~M)	0.72 (~D)	0.65 (~L)	0.68 (~C)	0.77 (ALL)	0.93 (~N)	0.81 (~D)	0.76 (C)	0.81 (~F)	0.9 (ALL)
Best buggy change recall	0.57 (ALL)	0.98 (M)	0.61 (~M)	0.63 (ALL)	0.62 (~L)	0.57 (~C)	0.57 (ALL)	0.63 (~F)	0.51 (~D)	0.43 (ALL)	0.8 (~D)	0.6 (~M)

### B. Important Individual Features

While the feature group importance provides some insight into the relative importance of groups of features, it does not say anything about individual features. Which features are the most important within a given project?

Using the chi-squared measure, all features are ranked individually. Additionally, the distribution of each feature in buggy and clean changes is computed to decide whether the corresponding feature contributes more to buggy or clean change classification. The top 5 ranked individual features in each feature group are determined, and presented in Table 9. Each box lists the top 5 ranked features within a feature group, for a given project. Each individual feature is listed, along with its overall numerical rank among the total set of features available for that project. The + and – before the rank indicate whether the feature is contributing to the buggy (+) or clean (-) change class.

**Table 9. Top five ranked individual features in each feature group.** Numbers in parentheses indicate the overall rank (computed using a Chi-square measure) of a feature’s importance among all features. A ‘+’ sign indicates the feature contributes to buggy changes, and a ‘–’ sign indicates the feature contributes to clean changes. The  $\Delta$  mark beside a complexity metric indicates it is a delta metric.

	Bugzilla	Eclipse	Plone
Complexity metrics	N/A	SumEssential(-117), $\Delta$ CountLineBlank(+228), CountStmtDecl(-417), CountLineComment(-419), CountLineCodeDecl(-420)	N/A
Change Log	fix(+345), comments (-351), correcting(-414), patch(+480), ability(- 492)	fix(+386), for(+398), 18(+961), 3249(+962), 1(- 1795)	action(+396), beautified(+576), catch(+577), categories(+578), global(+579)
Metadata	changed loc(+1), loc(+2), bug count(+3), time(-9), change count(+43)	time(+74), changed loc(+88), bug count(+104), days(-137), change log length(-142)	changed loc(+21), bug count(+140), loc(+274), time(-456), author(+676)
New Source	order(+4), b(+7), bit(+8), ok(+10), used(+11)	flowinfo(+1), analysecode(+2), flowcontext(+3), slow(+4), iabstractsynt(+6)	globals(-2), security(+3), not(+4), accesscontrol(+5), aq(+6)
Added Delta	if(+5), my(+6), value(+15), not(+22), sendsql(+26)	codestream(+12), recordpositionsfrom(+8), belongsto(+24), complete(+25), jobfamily(+26)	self(+1), def(+26), %(+40), raise(+46), log(+50)
Deleted Delta	name(+153), value(+281), my(+300), fetchsqldata(+326), sendsql(+349)	codestream(+14), public(+15), recordpositionsfrom(+19), return(+21), this(+22)	self(+14), %(+47), log(+55), def(+107), else(+143)
Directory/ File name	relation(-219), set(-220), move(-375), createattachment(- 490), export(-491)	ast(+5), compiler(+47), statement(+108), core (- 116), model(-214)	tool(+10), scripts(-154), edit(-340), form(-411), folder(+470)

For example, in the Bugzilla project in the Added Delta feature group, the keyword “if” is listed as the top feature within the group, and is the 5<sup>th</sup> most important individual feature overall. Compared to traditional bug prediction research that tends to use software metrics to determine

bugs in software, the most important individual features presented above seem to have limited utility for constructing causal models of what causes a bug. A general, cross-project causal model of bug injection just cannot explain why “if” is a strong feature for Bugzilla, while “self” is a strong feature for Plone. One explanation is that these are statistically correlated features computed for each project, and hence there should not be any expectation of a deeper model. The individual important features are deeply project-specific, indicating that no cross-project classification model can be developed. It is just not possible to train a classifier using these feature types on one project, then apply it to another project, and obtain any kind of reasonable accuracy, precision, or recall.

One interesting question is whether the committing developer is predictive for bugginess. Based on this data, the short answer is “no”, not for change classification. In Table 9 above, only one project, Plone, lists author as a top 5 feature within the metadata feature group, and it is ranked low, at 676. One explanation is that even bug-prone developers tend to make more clean changes than buggy changes. If, hypothetically, a developer made 85% clean changes and 15% buggy changes, they might be perceived by their peers to be a bug-prone developer, while the change classifier would view this same developer as a strong feature for predicting clean changes. An intriguing potential extension of change classification would be to train one change classifier for each developer on a project, and then perform developer-specific bug prediction. Developers are expected to have developer-specific patterns in their bug introducing changes, and this might improve the performance of change classification. This remains future work.



## VIII. DISCUSSION

This section discusses possible applications of change classification, and provides additional interpretation of the results. This section ends with some discussion on the limitations of our experiments.

### *A. Potential Applications of Change Classification*

Right now, the buggy change classifier operates in a lab environment. However, it could potentially be put into use in various ways:

- **A commit checker:** The classifier identifies buggy changes during commits of changes to a SCM system, and notifies developers of the results. The bug prediction in the commit checker is immediate, and makes it easy for developers to inspect the change just made.
- **Potential bug indicator during source code editing:** We have shown that features from source code (A, N, D) have discriminative power (see Figure 5). That is, just using features from source code, it is possible to perform accurate bug classification. This implies that a bug classifier can be embedded in a source code editor. During the source code editing process, the classifier could monitor source code changes. As soon as the cumulative set of changes made during an editing session leads the classifier to make a bug prediction, the editor can notify the developer. A proof of concept implementation of this idea using the Eclipse IDE is reported in [27].
- **Impact on the software development process:** Results from the change classifier could be integrated into the software development process. After committing a change, a developer receives feedback from the classifier. If the classifier indicates it was a buggy change, this could trigger an automatic code inspection on the change by multiple engineers. After the inspection, the developer commits a modified change and receives

more feedback. If this approach is effective in finding bugs right away, it could significantly reduce the number of latent bugs in a software system.

### *B. Issues of Change Classification*

A software project that records its changes in an SCM system, has approximately 100 or more SCM commits, and has some record of which commits fix bugs, may have the change classification technique used on it.

Several issues arise when extracting features from an existing project. First, an examination of change log messages is required to determine how best to determine the bug fix changes. The best set of keywords depends on how each project has used their SCM log messages in the past. For projects that consistently use a change tracking system, data may need to be extracted from this system as well.

Another concern is ensuring the memory capacity of the machine performing classification is not exceeded, as this leads to excessive swapping to disk, and reduced performance. Two techniques can be used to reduce memory usage. The range of revisions used to train the classifier can be reduced; this was the case in the current paper with the Eclipse and JEdit projects. The project can also be subdivided into smaller modules, such as considering only a specific subdirectory within a large project. A single, moderately high-end workstation computer should be sufficient for performing change classification work, and hence hardware costs for implementing the technique are modest.

An SVM classifier generally works well using all feature groups available for a project. If time is available, a given project can consider performing a feature group sensitivity analysis, of the type described in Section VII. This permits the use of the most accurate feature groups for the current project, usually resulting in a small gain in performance. Another consideration is


whether an existing complexity metric computation tool is available. Change classification can provide accurate results without the use of complexity metrics as a feature, however, complexity metrics are sometimes the most accurate single class of feature.

In an ongoing project, the SVM classifier will need to be periodically retrained to accommodate data from new project changes. If a project is small enough, training an SVM could be performed nightly, at the end of the workday. On larger projects, the SVM could be retrained weekly.

All in all, we expect that change classification will assist developers in identifying those changes that are most likely to contain bugs—and thus increase quality, reduce effort, or both.

### *C. Minimum Change Numbers for Classifier Training*

The results presented in this paper use changes in 500 (or 250 revisions) to train and evaluate an SVM classifier. This raises the question of how many revisions or changes are required to train an SVM classifier to yield reasonable classification performance. To answer this question, subsets of changes are used to train and evaluate an SVM classifier with all features. To begin, only the first 10 changes are used to train and evaluate a classifier. Next, the first 20 changes are used. In the same way, the number of changes is increased by 10 and then used to train and evaluate an SVM classifier.



**Figure 6.** Number of changes used to train and evaluate an SVM classifier and their corresponding accuracy of selected projects: Apache, Bugzilla, GForge, Mozilla, and SVN.

Figure 6 shows accuracies of selected projects by using various numbers of changes. Two conclusions can be drawn from this. First, after approximately 100 changes predictive accuracy is generally close to steady-state values. There is still some churn in accuracy from changes 100 to 200, but the accuracy does not have dramatic swings (for example, no  $\pm 20\%$  deviations) after this point. Accuracies settle down to steady-state values after 200 changes for most projects. It thus appears that change classification using an SVM classifier is usable as a bug prediction technique once a project has logged 100 changes, though with some variability in predictive accuracy for the next 100 changes until accuracy values reach steady-state. Due to this, if change classification is to be used on a new software project, it is probably best adopted towards the middle or end of the initial coding phase, once sufficient changes have been logged and initial testing efforts have begun to reveal project bugs. Any project that has already had at least one release would typically be able to adopt change classification at any time.

#### *D. Threats to Validity*

There are five major threats to the validity of this study.

**Systems examined might not be representative.** 12 systems are examined, more than any other work reported in the literature. In spite of this, it is still possible that we accidentally chose systems that have better (or worse) than average bug classification accuracy. Since we intentionally only chose systems that had some degree of linkage between change tracking systems and the text in the change log (so we could determine fix inducing changes), we have a project selection bias. This is most evident in the data from Bugzilla and Scarab, where the fact that they are change tracking systems led to a higher than normal ratio of buggy to clean changes.

**Systems are all open source.** The systems examined in this paper all use an open source development methodology, and hence might not be representative of all development contexts. It is possible that the stronger deadline pressure, different personnel turnover patterns, and different development processes used in commercial development could lead to different buggy change patterns.

**Bug fix data is incomplete.** Even though we selected projects that have change logs with good quality, we still are only able to extract a subset of the total number of bugs (typically only 40%-60% of those reported in the bug tracking system). Since the quality of change logs varies across projects, it is possible that the output of the classification algorithm will include false positive and false negatives. It is currently unclear what impact lower quality change logs has on the classification results.

**Bug introducing data is incomplete.** The SZZ algorithm used to identify bug-introducing changes has limitations: it cannot find bug introducing changes for bug fixes that only involve the deletion of source code. It also cannot identify bug-introducing changes caused by a change made to a file different from the one being analyzed. It is also possible to miss bug-introducing

changes when a file changes its name, since the algorithm does not track such name changes.

**Requires initial change data to train a classification model.** As discussed in Section C, the change classification technique requires about 100 changes to train a project specific classification model before predictive accuracy achieves a “usable” level of accuracy.

**Use of bug tracking systems for tracking new functionalities.** In two of the systems examined, Bugzilla and Scarab, the projects used bug tracking systems to also track new functionality additions to the project. For these projects, the meaning of a bug tracking identifier in the change log message either means a bug was fixed, or a new functionality added. This substantially increases the number of changes flagged as bug fixes. For these systems, the interpretation of a positive classification output is a change that is either buggy or a new functionality. When using this algorithm, care needs to be taken to understand the meaning of changes identified as bugs, and, wherever possible, to ensure that only truly buggy changes are flagged as being buggy.

## IX. CONCLUSION AND OPEN ISSUES

If a developer knows that a change she just made contains a bug, she can use this information to take steps to identify and fix the potential bug in the change before it leads to a bug report. This paper has introduced a new bug prediction technique that works at fine granularity, an individual file level change, and has accuracy comparable to the best existing bug prediction techniques in the literature (78%, on average). Features gathered only from source code have strong discriminative power, suggesting the possibility of embedding the classification algorithm into integrated development environments for bug prediction during editing sessions. Developers can benefit from focused and prompt prediction of buggy changes, receiving this prediction either while they are editing source code or right after a change submission.

This work is the first to classify file changes as buggy or clean using the combination of change information features and source code terms. Additionally, this work provides an evaluation of the relative contributions of various feature groups for change classification.

Although these experimental results are encouraging, there are still several open issues in this work, including:

- Exploring on-line machine learning algorithms to learn and update a classification model as the project progresses and using it to predict future changes.
- Generating more features from change information and exploring various ways to extract features, such as latent semantic analysis [22].
- Deep analysis on the individual features to identify common bug-prone code patterns or causality of bugs.
- Applying or modifying existing machine learning algorithms to achieve better prediction accuracy, precision, and recall [9].
- Finally, use in ongoing software projects to see how developers exploit bug classification information to perform bug reducing interventions in their software.

Overall, we expect that future approaches will see software history not only as a series of revisions and changes, but also as a series of successes and failures—and as a source for continuous awareness and improvement. Change classification is a first step in this direction.

#### ACKNOWLEDGMENT

We thank David P. Helmbold for his guidance and insight on the machine learning algorithms for file change classifications. We thank Kai Pan and Mark Slater for their valuable feedback on this paper. We especially thank Kevin Greenan and the Storage Systems Research Center at UCSC for allowing the use of their cluster for our research. Comments by the anonymous

reviewers were very helpful in improving the paper, and are greatly appreciated.

## REFERENCES

- [1] E. Alpaydin, *Introduction to Machine Learning*: The MIT Press, 2004.
- [2] G. Antoniol, G. Casazza, and A. Cimitile, "Traceability Recovery by Modelling Programmer Behaviour," *Proc. of 7th Working Conference on Reverse Engineering*, Brisbane, Queensland, Australia, pp. 240-247., 2000.
- [3] J. Anvik, L. Hiew, and G. C. Murphy, "Who Should Fix This Bug?" *Proc. of 2006 International Conference on Software Engineering (ICSE 2006)*, Shanghai, China, pp. 361-370, 2006.
- [4] J. Bevan, E. J. Whitehead, Jr., S. Kim, and M. Godfrey, "Facilitating Software Evolution with Kenyon," *Proc. of the 2005 European Software Engineering Conference and 2005 Foundations of Software Engineering (ESEC/FSE 2005)*, Lisbon, Portugal, pp. 177-186, 2005.
- [5] Y. Brun and M. D. Ernst, "Finding Latent Code Errors via Machine Learning over Program Executions," *Proc. of 26th International Conference on Software Engineering (ICSE 2004)*, Scotland, UK, pp. 480-490, 2004.
- [6] D. Cubranic and G. C. Murphy, "Hipikat: Recommending pertinent software development artifacts," *Proc. of 25th International Conference on Software Engineering (ICSE 2003)*, Portland, Oregon, pp. 408-418, 2003.
- [7] M. Fischer, M. Pinzger, and H. Gall, "Populating a Release History Database from Version Control and Bug Tracking Systems," *Proc. of 19th International Conference on Software Maintenance (ICSM 2003)*, Amsterdam, The Netherlands, pp. 23-32, 2003.
- [8] C. Flanagan, K. R. M. Leino, M. Lillibridge, G. Nelson, J. B. Saxe, and R. Stata, "Extended Static Checking for Java," *Proc. of the ACM SIGPLAN 2002 Conference on Programming Language Design and Implementation*, Berlin, Germany, pp. 234-245, 2002.
- [9] Y. Freund and R. E. Schapire, "A Short Introduction to Boosting," *Journal of Japanese Society for Artificial Intelligence*, vol. 14, no. 5, pp. 771-780, 1999.
- [10] G. A. Di Lucca, M. Di Penta, and S. Gradara, "An Approach to Classify Software Maintenance Requests," *Proc. of 18th International Conference on Software Maintenance (ICSM 2002)*, Montréal, Canada, pp. 93-102, 2002.
- [11] T. L. Graves, A. F. Karr, J. S. Marron, and H. Siy, "Predicting Fault Incidence Using Software Change History," *IEEE Transactions on Software Engineering*, vol. 26, no. 7, pp. 653-661, 2000.
- [12] T. Gyimothy, R. Ferenc, and I. Siket, "Empirical Validation of Object-Oriented Metrics on Open Source Software for Fault Prediction," *IEEE Transactions on Software Engineering*, vol. 31, no. 10, pp. 897-910, 2005.
- [13] A. E. Hassan and R. C. Holt, "The Top Ten List: Dynamic Fault Prediction," *Proc. of 21st International Conference on Software Maintenance (ICSM 2005)*, Budapest, Hungary, pp. 263-272, 2005.
- [14] T. Joachims, "Text Categorization with Support Vector Machines: Learning with Many Relevant Features," *Proc. of ECML 98, 10th European Conference on Machine Learning*, Chemnitz, Germany, pp. 137-142, 1998.
- [15] T. M. Khoshgoftaar and E. B. Allen, "Ordering Fault-Prone Software Modules," *Software Quality Control Journal*, vol. 11, no. 1, pp. 19-37, 2003.
- [16] T. M. Khoshgoftaar and E. B. Allen, "Predicting the Order of Fault-Prone Modules in Legacy Software," *Proc. of The Ninth International Symposium on Software Reliability Engineering*, Paderborn, Germany, pp. 344-353, 1998.
- [17] S. Kim, K. Pan, and E. J. Whitehead, Jr., "Memories of Bug Fixes," *Proc. of the 2006 ACM SIGSOFT Foundations of Software Engineering (FSE 2006)*, Portland, Oregon, pp. 35 - 45, 2006.
- [18] S. Kim, T. Zimmermann, E. J. Whitehead, Jr., and A. Zeller, "Predicting Bugs from Cached History," *Proc. of the 29th International Conference on Software Engineering (ICSE 2007)*, Minneapolis, USA, pp. 489-498, 2007.
- [19] R. Krovetz, S. Ugurel, and C. L. Giles, "Classification of Source Code Archives," *Proc. of the 26th International ACM SIGIR Conference on Research and Development in Information Retrieval*, Toronto, Canada, pp. 425-426, 2003.
- [20] A. Kuhn, S. Ducasse, and T. Girba, "Enriching Reverse Engineering with Semantic Clustering," *Proc. of 12th Working Conference on Reverse Engineering (WCRE 2005)*, Pittsburgh, Pennsylvania, USA, pp. 133-142, 2005.
- [21] R. Kumar, S. Rai, and J. L. Trahan, "Neural-Network Techniques for Software-quality Evaluation," *Proc. of 1998 Reliability and Maintainability Symposium*, Anaheim, CA, pp. 155-161, 1998.
- [22] T. K. Landauer, P. W. Foltz, and D. Laham, "Introduction to Latent Semantic Analysis," *Discourse Processes*, vol. 25, pp. 259-284, 1998.
- [23] D. Lewis, Y. Yang, T. Rose, and F. Li, "RCV1: A New Benchmark Collection for Text Categorization Research," *Journal of Machine Learning Research*, vol. 5, pp. 361-397, 2004.
- [24] Z. Li and Y. Zhou, "PR-Miner: Automatically Extracting Implicit Programming Rules and Detecting Violations in Large Software Code," *Proc. of the 2005 European Software Engineering Conference and 2005 Foundations of Software Engineering (ESEC/FSE 2005)*, Lisbon, Portugal, pp. 306 - 315, 2005.
- [25] B. Livshits and T. Zimmermann, "DynaMine: Finding Common Error Patterns by Mining Software Revision Histories," *Proc. of the 2005 European Software Engineering Conference and 2005 Foundations of Software Engineering (ESEC/FSE 2005)*, Lisbon, Portugal, pp. 296-305, 2005.
- [26] J. Lyle and M. Weiser, "Automatic Program Bug Location by Program Slicing," *Proc. of 2nd International Conference on Computers and Applications*, pp. 877-883, 1987.
- [27] J. Madhavan and E. J. Whitehead, Jr., "Predicting Changes Inside an Integrated Development Environment," *Proc. of 2007 Eclipse Technology Exchange Workshop*, Montreal, Canada, 2007.
- [28] J. I. Maletic and N. Valluri, "Automatic Software Clustering via Latent Semantic Analysis," *Proc. of 14th IEEE International Conference on Automated Software Engineering (ASE'99)*, Cocoa Beach, Florida, USA, pp. 251, 1999.
- [29] A. Marcus and J. I. Maletic, "Recovering Documentation-to-Source-Code Traceability Links using Latent Semantic Indexing," *Proc. of the 25th International Conference on Software Engineering*, Portland, Oregon, pp. 125-135, 2003.
- [30] T. Menzies, J. Greenwald, and A. Frank, "Data Mining Static Code Attributes to Learn Defect Predictors," *IEEE Trans. Software Engineering*, vol. 33, no. 1, pp. 2-13, 2007.



- [31] O. Mizuno and T. Kikuno, "Training on Errors Experiment to Detect Fault-prone Software Modules by Spam Filter," *Proc. of the 2007 European Software Engineering Conference and 2007 Foundations of Software Engineering (ESEC/FSE 2007)*, Dubrovnik, Croatia, pp. 405 - 414, 2007.
- [32] A. Mockus and L. G. Votta, "Identifying Reasons for Software Changes Using Historic Databases," *Proc. of 16th International Conference on Software Maintenance (ICSM 2000)*, San Jose, California, USA, pp. 120-130, 2000.
- [33] A. Mockus and D. M. Weiss, "Predicting Risk of Software Changes," *Bell Labs Technical Journal*, vol. 5, no. 2, pp. 169-180, 2002.
- [34] D. C. Montgomery, G. C. Runger, and N. F. Hubele, *Engineering Statistics*: Wiley, 2001.
- [35] A. W. Moore, "Cross-Validation," 2005, <http://www.autonlab.org/tutorials/overfit.html>.
- [36] N. Nagappan and T. Ball, "Use of Relative Code Churn Measures to Predict System Defect Density," *Proc. of 27th International Conference on Software Engineering (ICSE 2005)*, Saint Louis, Missouri, USA, pp. 284-292, 2005.
- [37] D. J. Newman, S. Hettich, C. L. Blake, and C. J. Merz, "UCI Repository of machine learning databases," 1988, <http://www.ics.uci.edu/~mllearn/MLRepository.html>.
- [38] T. J. Ostrand and E. J. Weyuker, "The Distribution of Faults in a Large Industrial Software System," *Proc. of the 2002 ACM SIGSOFT International Symposium on Software Testing and Analysis*, Roma, Italy, pp. 55-64, 2002.
- [39] T. J. Ostrand, E. J. Weyuker, and R. M. Bell, "Predicting the Location and Number of Faults in Large Software Systems," *IEEE Transactions on Software Engineering*, vol. 31, no. 4, pp. 340-355, 2005.
- [40] T. J. Ostrand, E. J. Weyuker, and R. M. Bell, "Where the Bugs Are," *Proc. of 2004 ACM SIGSOFT International Symposium on Software Testing and Analysis*, Boston, Massachusetts, USA, pp. 86-96, 2004.
- [41] K. Pan, S. Kim, and J. E. James Whitehead, "Bug Classification Using Program Slicing Metrics," *Proc. of the Sixth IEEE International Workshop on Source Code Analysis and Manipulation (SCAM 2006)*, Philadelphia, PA, 2006.
- [42] M. D. Penta, S. Gradara, and G. Antoniol, "Traceability Recovery in RAD Software Systems," *Proc. of the 10th IEEE International Workshop on Program Comprehension*, Paris, France, pp. 207-216., 2002.
- [43] B. Raskutti, H. L. Ferrá, and A. Kowalczyk, "Second Order Features for Maximizing Text Classification Performance," *Proc. of 12th European Conference on Machine Learning*, Freiburg, Germany, pp. 419-430, 2001.
- [44] Scientific Toolworks, "Maintenance, Understanding, Metrics and Documentation Tools for Ada, C, C++, Java, and FORTRAN," 2005, <http://www.scitools.com/>.
- [45] S. Scott and S. Matwin, "Feature Engineering for Text Classification," *Proc. of the Sixteenth International Conference on Machine Learning*, Bled, Slovenia, pp. 379-388, 1999.
- [46] F. Sebastiani, "Machine Learning in Automated Text Categorization," *ACM Computing Surveys*, vol. 34, no. 1, pp. 1-47, 2002.
- [47] J. Śliwerski, T. Zimmermann, and A. Zeller, "When Do Changes Induce Fixes?" *Proc. of Int'l Workshop on Mining Software Repositories (MSR 2005)*, Saint Louis, Missouri, USA, pp. 24-28, 2005.
- [48] J. C. Spohrer, E. Soloway, and E. Pope, "Where the Bugs Are," *Proc. of the SIGCHI Conference on Human Factors in Computing Systems*, San Francisco, California, USA, pp. 47-53, 1985.
- [49] V. N. Vapnik, *The Nature of Statistical Learning Theory*: Springer-Verlag, 1995.
- [50] Wikipedia, "CamelCase," 2005, <http://en.wikipedia.org/wiki/CamelCase>.
- [51] C. C. Williams and J. K. Hollingsworth, "Automatic Mining of Source Code Repositories to Improve Bug Finding Techniques," *IEEE Trans. Software Engineering*, vol. 31, no. 6, pp. 466-480, 2005.
- [52] I. H. Witten and E. Frank, *Data Mining: Practical Machine Learning Tools and Techniques (Second Edition)*: Morgan Kaufmann, 2005.
- [53] L. Zhang, J. Zhu, and T. Yao, "An Evaluation of Statistical Spam Filtering Techniques," *ACM Transactions on Asian Language Information Processing (TALIP)*, vol. 3, no. 4, pp. 243-269, 2004.
- [54] Z. Zheng, X. Wu, and R. Srihari, "Feature selection for text categorization on imbalanced data," *ACM SIGKDD Explorations Newsletter*, vol. 6, no. 1, pp. 80-89, 2004.
- [55] T. Zimmermann and P. Weißgerber, "Preprocessing CVS Data for Fine-Grained Analysis," *Proc. of Int'l Workshop on Mining Software Repositories (MSR 2004)*, Edinburgh, Scotland, pp. 2-6, 2004.
- [56] T. Zimmermann, P. Weißgerber, S. Diehl, and A. Zeller, "Mining Version Histories to Guide Software Changes," *IEEE Trans. Software Engineering*, vol. 31, no. 6, pp. 429-445, 2005.

**Sunghun Kim** is a postdoctoral associate at the Massachusetts Institute of Technology. He received his PhD in Computer Science from the University of California, Santa Cruz in 2006. His main research interests are dynamic bug prediction and identifying common bug patterns by mining software repositories.

**E. James Whitehead, Jr.** (M'93) is an Associate Professor of Computer Science at the University of California, Santa Cruz. He received his PhD in Information and Computer Science in 2000 from the University of California, Irvine, and his BS in Electrical Engineering from the Rensselaer Polytechnic Institute in 1989. His research interests include software evolution, software design, application layer network protocols, collaborative authoring, computer game design and hypertext systems.

**Yi Zhang** (M'06) is an Assistant Professor at the University of California Santa Cruz. She received her Ph.D. and M.S. from Carnegie Mellon University. She was born in 1976 in Chengdu, China.